

In [20]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

In [21]:

```
train_df=pd.read_csv(r"C:\Users\DHEEPAK\Desktop\Mobile_Price_Classification_train.csv")
train_df
```

Out[21]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n
0	842	0	2.2	0	1	0	7	0.6	188	
1	1021	1	0.5	1	0	1	53	0.7	136	
2	563	1	0.5	1	2	1	41	0.9	145	
3	615	1	2.5	0	0	0	10	0.8	131	
4	1821	1	1.2	0	13	1	44	0.6	141	
...
1995	794	1	0.5	1	0	1	2	0.8	106	
1996	1965	1	2.6	1	0	0	39	0.2	187	
1997	1911	0	0.9	1	1	1	36	0.7	108	
1998	1512	0	0.9	0	4	1	46	0.1	145	
1999	510	1	2.0	1	5	1	45	0.9	168	

2000 rows × 21 columns



In [22]:

```
test_df=pd.read_csv(r"C:\Users\DHEEPAK\Desktop\Mobile_Price_Classification_test.csv")
test_df
```

Out[22]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_v
0	1	1043	1	1.8	1	14	0	5	0.1	15
1	2	841	1	0.5	1	4	1	61	0.8	15
2	3	1807	1	2.8	0	1	0	27	0.9	18
3	4	1546	0	0.5	1	18	1	25	0.5	9
4	5	1434	0	1.4	0	11	1	49	0.5	10
...
995	996	1700	1	1.9	0	0	1	54	0.5	17
996	997	609	0	1.8	1	0	0	13	0.9	18
997	998	1185	0	1.4	0	1	1	8	0.5	8
998	999	1533	1	0.5	1	0	0	50	0.4	17
999	1000	1270	1	0.5	0	4	1	35	0.1	14

1000 rows × 21 columns



In [23]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   battery_power          2000 non-null   int64
1   blue                   2000 non-null   int64
2   clock_speed            2000 non-null   float64
3   dual_sim               2000 non-null   int64
4   fc                     2000 non-null   int64
5   four_g                 2000 non-null   int64
6   int_memory             2000 non-null   int64
7   m_dep                  2000 non-null   float64
8   mobile_wt              2000 non-null   int64
9   n_cores                2000 non-null   int64
10  pc                     2000 non-null   int64
11  px_height              2000 non-null   int64
12  px_width               2000 non-null   int64
13  ram                    2000 non-null   int64
14  sc_h                   2000 non-null   int64
15  sc_w                   2000 non-null   int64
16  talk_time              2000 non-null   int64
17  three_g                2000 non-null   int64
18  touch_screen           2000 non-null   int64
19  wifi                   2000 non-null   int64
20  price_range            2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

In [24]:

```
test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   id                  1000 non-null   int64
 1   battery_power       1000 non-null   int64
 2   blue                 1000 non-null   int64
 3   clock_speed         1000 non-null   float64
 4   dual_sim            1000 non-null   int64
 5   fc                  1000 non-null   int64
 6   four_g              1000 non-null   int64
 7   int_memory          1000 non-null   int64
 8   m_dep               1000 non-null   float64
 9   mobile_wt           1000 non-null   int64
10   n_cores             1000 non-null   int64
11   pc                  1000 non-null   int64
12   px_height           1000 non-null   int64
13   px_width            1000 non-null   int64
14   ram                 1000 non-null   int64
15   sc_h                1000 non-null   int64
16   sc_w                1000 non-null   int64
17   talk_time           1000 non-null   int64
18   three_g             1000 non-null   int64
19   touch_screen        1000 non-null   int64
20   wifi                1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

In [25]:

```
x=train_df.drop('dual_sim',axis=1)
y=train_df['dual_sim']
```

In [26]:

```
x=test_df.drop('dual_sim',axis=1)
y=test_df['dual_sim']
```

In [27]:

```
train_df['blue'].value_counts()
```

Out[27]:

```
blue
0    1010
1     990
Name: count, dtype: int64
```

In [28]:

```
test_df['blue'].value_counts()
```

Out[28]:

blue
1 516
0 484
Name: count, dtype: int64

In [29]:

```
T={"three_g":{'Yes':1,'No':0}}  
train_df=train_df.replace(T)  
print(train_df)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory			
0	842	0	2.2	0	1	0	7	\		
1	1021	1	0.5	1	0	1	53			
2	563	1	0.5	1	2	1	41			
3	615	1	2.5	0	0	0	10			
4	1821	1	1.2	0	13	1	44			
...			
1995	794	1	0.5	1	0	1	2			
1996	1965	1	2.6	1	0	0	39			
1997	1911	0	0.9	1	1	1	36			
1998	1512	0	0.9	0	4	1	46			
1999	510	1	2.0	1	5	1	45			
	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	
0	0.6	188	2	...	20	756	2549	9	7	\
1	0.7	136	3	...	905	1988	2631	17	3	
2	0.9	145	5	...	1263	1716	2603	11	2	
3	0.8	131	6	...	1216	1786	2769	16	8	
4	0.6	141	2	...	1208	1212	1411	8	2	
...	
1995	0.8	106	6	...	1222	1890	668	13	4	
1996	0.2	187	4	...	915	1965	2032	11	10	
1997	0.7	108	8	...	868	1632	3057	9	1	
1998	0.1	145	5	...	336	670	869	18	10	
1999	0.9	168	6	...	483	754	3919	19	4	
	talk_time	three_g	touch_screen	wifi	price_range					
0	19	0	0	1	1					
1	7	1	1	0	2					
2	9	1	1	0	2					
3	11	1	0	0	2					
4	15	1	1	0	1					
...					
1995	19	1	1	0	0					
1996	16	1	1	1	2					
1997	5	1	1	0	3					
1998	19	1	1	1	0					
1999	2	1	1	1	3					

[2000 rows x 21 columns]

In []:

1

In [30]:

```
T={"three_g":{"Yes":1,'No':0}}
test_df=test_df.replace(T)
print(test_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	1	1043	1	1.8	1	14	0	5
\								
1	2	841	1	0.5	1	4	1	61
2	3	1807	1	2.8	0	1	0	27
3	4	1546	0	0.5	1	18	1	25
4	5	1434	0	1.4	0	11	1	49
..
995	996	1700	1	1.9	0	0	1	54
996	997	609	0	1.8	1	0	0	13
997	998	1185	0	1.4	0	1	1	8
998	999	1533	1	0.5	1	0	0	50
999	1000	1270	1	0.5	0	4	1	35

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	\
0	0.1	193	...	16	226	1412	3476	12	7	
1	0.8	191	...	12	746	857	3895	6	0	
2	0.9	186	...	4	1270	1366	2396	17	10	
3	0.5	96	...	20	295	1752	3893	10	0	
4	0.5	108	...	18	749	810	1773	15	8	
..	
995	0.5	170	...	17	644	913	2121	14	8	
996	0.9	186	...	2	1152	1632	1933	8	1	
997	0.5	80	...	12	477	825	1223	5	0	
998	0.4	171	...	12	38	832	2509	15	11	
999	0.1	140	...	19	457	608	2828	9	2	

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

In [31]:

```
x=train_df.drop('dual_sim',axis=1)
y=train_df['dual_sim']
```

In [32]:

```
x=test_df.drop('dual_sim',axis=1)
y=test_df['dual_sim']
```

In [33]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

Out[33]:

```
((700, 20), (300, 20))
```

In [34]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[34]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [35]:

```
rf=RandomForestClassifier()
```

In [36]:

```
params={'max_depth':[2,3,5,10,20],
        'min_samples_leaf':[5,10,20,50,100,200],
        'n_estimators':[10,25,30,50,100,200]}
```

In [44]:

```
from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_train,y_train)
```

Out[44]:

```
► GridSearchCV
► estimator: RandomForestClassifier
  ► RandomForestClassifier
```

In [40]:

```
grid_search.best_score_
```

Out[40]:

```
0.5557142857142857
```

In [46]:

```
rf_best=grid_search.best_estimator_  
print(rf_best)
```

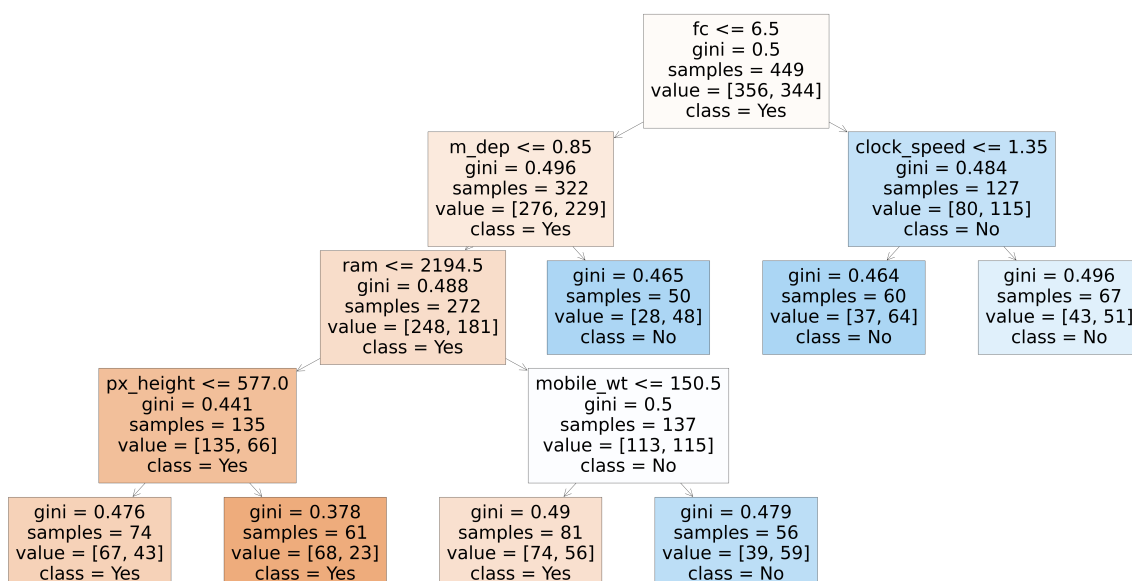
RandomForestClassifier(max_depth=20, min_samples_leaf=50, n_estimators=25)

In [50]:

```
from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))  
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

Out[50]:

```
[Text(0.6363636363636364, 0.9, 'fc <= 6.5\ngini = 0.5\nsamples = 449\nvalue =  
[356, 344]\nnclass = Yes'),  
Text(0.45454545454545453, 0.7, 'm_dep <= 0.85\ngini = 0.496\nsamples = 322\nvalue = [276, 229]\nnclass = Yes'),  
Text(0.36363636363636365, 0.5, 'ram <= 2194.5\ngini = 0.488\nsamples = 272\nvalue = [248, 181]\nnclass = Yes'),  
Text(0.18181818181818182, 0.3, 'px_height <= 577.0\ngini = 0.441\nsamples = 135\nvalue = [135, 66]\nnclass = Yes'),  
Text(0.09090909090909091, 0.1, 'gini = 0.476\nsamples = 74\nvalue = [67, 43]\nnclass = Yes'),  
Text(0.2727272727272727, 0.1, 'gini = 0.378\nsamples = 61\nvalue = [68, 23]\nnclass = Yes'),  
Text(0.5454545454545454, 0.3, 'mobile_wt <= 150.5\ngini = 0.5\nsamples = 137\nvalue = [113, 115]\nnclass = No'),  
Text(0.45454545454545453, 0.1, 'gini = 0.49\nsamples = 81\nvalue = [74, 56]\nnclass = Yes'),  
Text(0.6363636363636364, 0.1, 'gini = 0.479\nsamples = 56\nvalue = [39, 59]\nnclass = No'),  
Text(0.5454545454545454, 0.5, 'gini = 0.465\nsamples = 50\nvalue = [28, 48]\nnclass = No'),  
Text(0.8181818181818182, 0.7, 'clock_speed <= 1.35\ngini = 0.484\nsamples = 127\nvalue = [80, 115]\nnclass = No'),  
Text(0.7272727272727273, 0.5, 'gini = 0.464\nsamples = 60\nvalue = [37, 64]\nnclass = No'),  
Text(0.9090909090909091, 0.5, 'gini = 0.496\nsamples = 67\nvalue = [43, 51]\nnclass = No')]
```

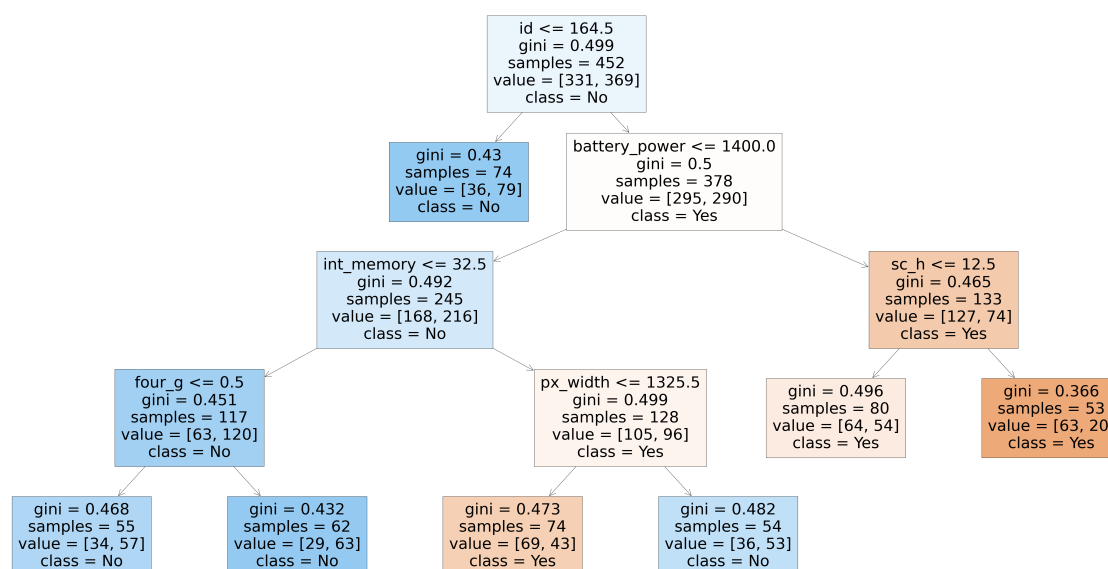


In [51]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

Out[51]:

```
[Text(0.5, 0.9, 'id <= 164.5\ngini = 0.499\nsamples = 452\nvalue = [331, 369]\n\nclass = No'),
Text(0.4090909090909091, 0.7, 'gini = 0.43\nsamples = 74\nvalue = [36, 79]\n\nclass = No'),
Text(0.5909090909090909, 0.7, 'battery_power <= 1400.0\ngini = 0.5\nsamples = 378\nvalue = [295, 290]\n\nclass = Yes'),
Text(0.36363636363636365, 0.5, 'int_memory <= 32.5\ngini = 0.492\nsamples = 245\nvalue = [168, 216]\n\nclass = No'),
Text(0.18181818181818182, 0.3, 'four_g <= 0.5\ngini = 0.451\nsamples = 117\nvalue = [63, 120]\n\nclass = No'),
Text(0.09090909090909091, 0.1, 'gini = 0.468\nsamples = 55\nvalue = [34, 57]\n\nclass = No'),
Text(0.2727272727272727, 0.1, 'gini = 0.432\nsamples = 62\nvalue = [29, 63]\n\nclass = No'),
Text(0.5454545454545454, 0.3, 'px_width <= 1325.5\ngini = 0.499\nsamples = 128\nvalue = [105, 96]\n\nclass = Yes'),
Text(0.45454545454545453, 0.1, 'gini = 0.473\nsamples = 74\nvalue = [69, 43]\n\nclass = Yes'),
Text(0.6363636363636364, 0.1, 'gini = 0.482\nsamples = 54\nvalue = [36, 53]\n\nclass = No'),
Text(0.8181818181818182, 0.5, 'sc_h <= 12.5\ngini = 0.465\nsamples = 133\nvalue = [127, 74]\n\nclass = Yes'),
Text(0.7272727272727273, 0.3, 'gini = 0.496\nsamples = 80\nvalue = [64, 54]\n\nclass = Yes'),
Text(0.9090909090909091, 0.3, 'gini = 0.366\nsamples = 53\nvalue = [63, 20]\n\nclass = Yes')]
```



In [52]:

```
rf_best.feature_importances_
```

Out[52]:

```
array([0.10818799, 0.11620779, 0.00616694, 0.0746391 , 0.08398264,  
       0.00038537, 0.02412626, 0.04710012, 0.05134397, 0.02683694,  
       0.09601713, 0.07501672, 0.05224792, 0.06330306, 0.04853793,  
       0.08064983, 0.00584961, 0.          , 0.02331212, 0.01608855])
```

In [54]:

```
imp_df=pd.DataFrame({'Varname':x_train.columns,"Imp":rf_best.feature_importances_})  
imp_df.sort_values(by="Imp",ascending=False)
```

Out[54]:

	Varname	Imp
1	battery_power	0.116208
0	id	0.108188
10	pc	0.096017
4	fc	0.083983
15	sc_w	0.080650
11	px_height	0.075017
3	clock_speed	0.074639
13	ram	0.063303
12	px_width	0.052248
8	mobile_wt	0.051344
14	sc_h	0.048538
7	m_dep	0.047100
9	n_cores	0.026837
6	int_memory	0.024126
18	touch_screen	0.023312
19	wifi	0.016089
2	blue	0.006167
16	talk_time	0.005850
5	four_g	0.000385
17	three_g	0.000000

In []:

